

OB DESCRIPTION

Generative AI Engineer – LLM, RAG & AI Applications

XYZ AI Solutions | Bengaluru / Hybrid | Full-Time

ABOUT THE ROLE

We are looking for a Generative AI Engineer to design, develop, evaluate, and deploy reliable AI-powered applications using Python, Large Language Models, Retrieval-Augmented Generation, APIs, cloud platforms, and modern software engineering practices. The role involves working with engineering, product, and data teams to turn business requirements into practical AI solutions.

KEY RESPONSIBILITIES

- Build and maintain Generative AI and LLM applications using Python.
- Design RAG pipelines using embeddings, vector databases, semantic search, and retrieval strategies.
- Develop backend services and REST APIs using FastAPI, Flask, or equivalent Python frameworks.
- Design prompts, structured outputs, evaluation workflows, and basic guardrails for LLM applications.
- Integrate AI applications with internal systems, third-party APIs, and enterprise data sources.
- Build data-processing, ETL, and data-pipeline workflows for AI and analytics workloads.
- Containerize Python services using Docker and support deployment on cloud infrastructure.
- Work with Git, GitHub, CI/CD workflows, and software engineering best practices.
- Evaluate AI response quality, relevance, and reliability using representative test cases and quality metrics.
- Collaborate with engineering, product, and data teams to deliver measurable business outcomes.

REQUIRED SKILLS

- Strong Python programming and software engineering fundamentals.
- Hands-on experience with Generative AI and Large Language Models (LLMs).
- Practical experience building Retrieval-Augmented Generation (RAG) applications.
- Experience with embeddings, vector databases, and semantic search.
- Prompt engineering and LLM/AI evaluation.
- REST API development using FastAPI, Flask, or equivalent.
- SQL and database fundamentals.
- Data processing, ETL, or data pipelines.

- Docker and containerization.
- Git and GitHub or equivalent version control.
- CI/CD fundamentals.
- Cloud experience with AWS, Azure, or GCP.

PREFERRED SKILLS

- LangChain, LlamaIndex, or LangGraph.
- Hugging Face Transformers.
- AI agents or agentic workflows.
- Kubernetes or another container orchestration platform.
- Infrastructure as Code such as Terraform.
- Event-driven systems such as Apache Kafka.
- Workflow orchestration such as Airflow.
- Caching technologies such as Redis.
- Observability using Prometheus, Grafana, or OpenTelemetry.
- Automated LLM evaluation and AI quality metrics.

EXPERIENCE

2–5 years of experience in Python development, AI/ML engineering, Generative AI, backend engineering, data engineering, or a closely related technical field.

EDUCATION

Bachelor's or Master's degree in Computer Science, Information Technology, Data Science, Artificial Intelligence, Engineering, or a related discipline.

WHAT WE VALUE

- Strong analytical and problem-solving skills.
- Ability to explain technical concepts clearly.
- Ownership and ability to work independently.
- Collaborative approach to cross-functional engineering.
- Focus on production quality, measurable outcomes, and continuous improvement.

TECHNOLOGY ENVIRONMENT

Python • FastAPI • Flask • REST APIs • SQL • PostgreSQL • AWS • Docker • Git • GitHub • GitHub Actions • CI/CD • LangChain • LlamaIndex • Hugging Face • RAG • Vector Databases • Embeddings • Semantic Search • LLMs • Generative AI • AI Agents • Kubernetes • Prometheus • Grafana

EQUAL OPPORTUNITY

We evaluate candidates based on relevant skills, experience, problem-solving ability, and potential to contribute to the role.